

# Concurrent Programming: Threads

CS 18000

Sunil Prabhakar

Department of Computer Science

Purdue University



# [Objectives]

- This week we will get introduced to concurrent programming
  - Creating a new thread of execution
  - Waiting for threads to terminate
  - Thread states and scheduling
  - `sleep()` and `yield()`
  - Simple synchronization among threads

# [One-track mind?]

- Often, in real life we perform multiple tasks at the same time
  - Doing the laundry
  - Making a pot of coffee
- This is more efficient.
- Our programs thus far have had a single track (thread) of execution
  - at any point in time, there is only one statement being executed
  - not always efficient -- can stall (e.g., user input)

# [ Multiple Concurrent Threads ]

- There are many instances in computing where we can benefit from multiple concurrent threads of execution.
- For example:
  - GUI responsiveness. The GUI should not freeze while performing time-consuming operations.
  - Liveness in games: display shouldn't "lock up"
  - Exploiting available processing: speeding up processing by using all computing cores

# [ Motivation: Gaming ]

- Consider a game program that has to repeatedly
  - redraw the scene
  - play the game, record scores, ask the user if they want to play again.
- We don't want to stop redrawing the scene while waiting for the user input.
- *Solution: perform both tasks at the same time (concurrently)*

# [ Motivation: GUIs ]

- Consider a GUI event which causes some time-consuming processing to execute.
- While this processing is going on, the GUI will "lock up"
- This is not desirable from a user's experience point of view.
- How can we prevent this?
- *Solution: perform non-GUI processing without locking up GUI thread*

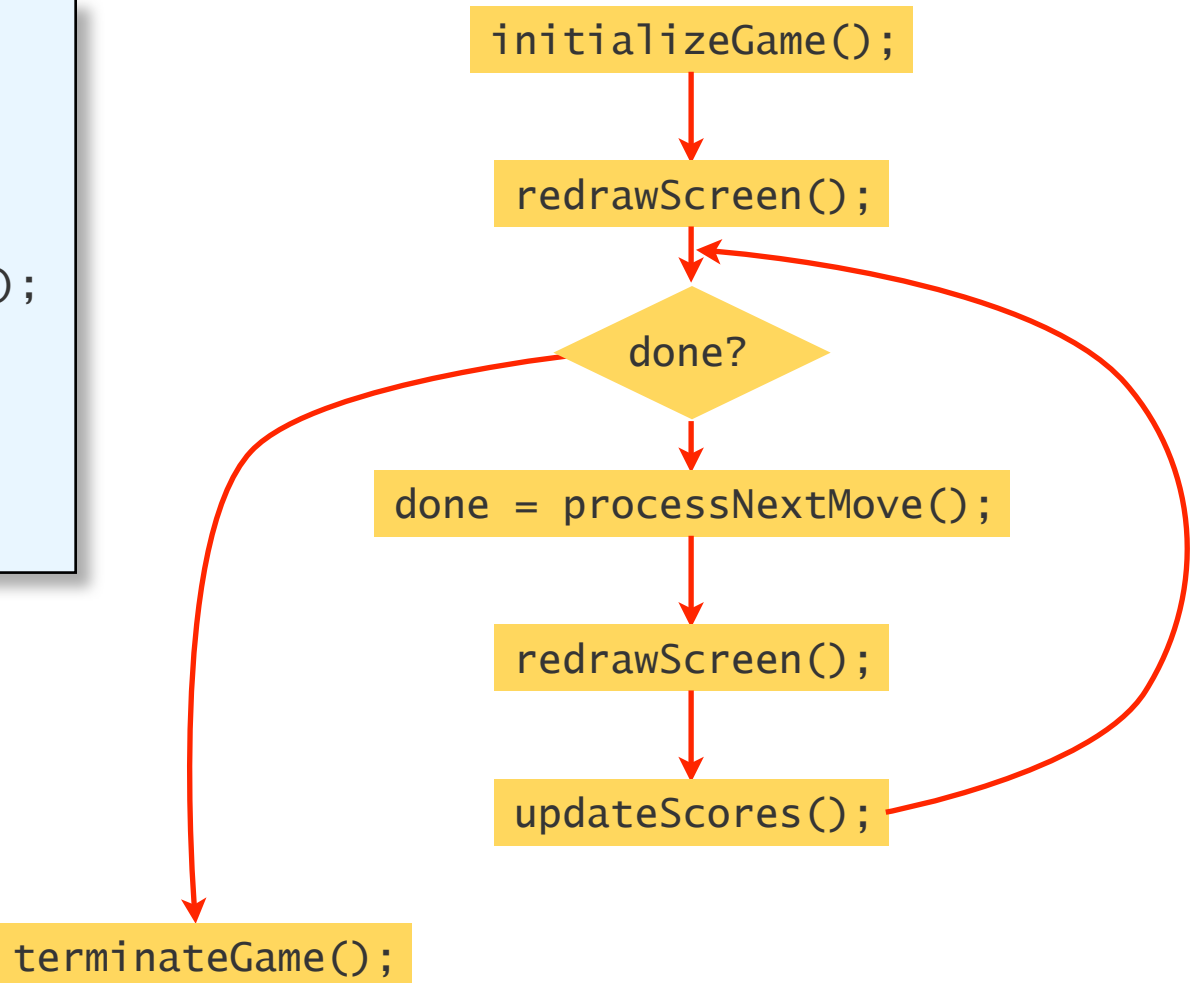
# [ Motivation: Exploiting Multiple Cores and Processors ]

- Due to the recent hardware trends, modern computers have multiple CPUs (cores or processors)
- If there is only a single thread of execution, only one CPU is used for our program.
- How do we exploit these other CPUs?
- Consider
  - the initialization of a large array
  - searching for an item in a large array
- *Solution: Split array into pieces and initialize (search) each piece concurrently.*

# Game: sequential version

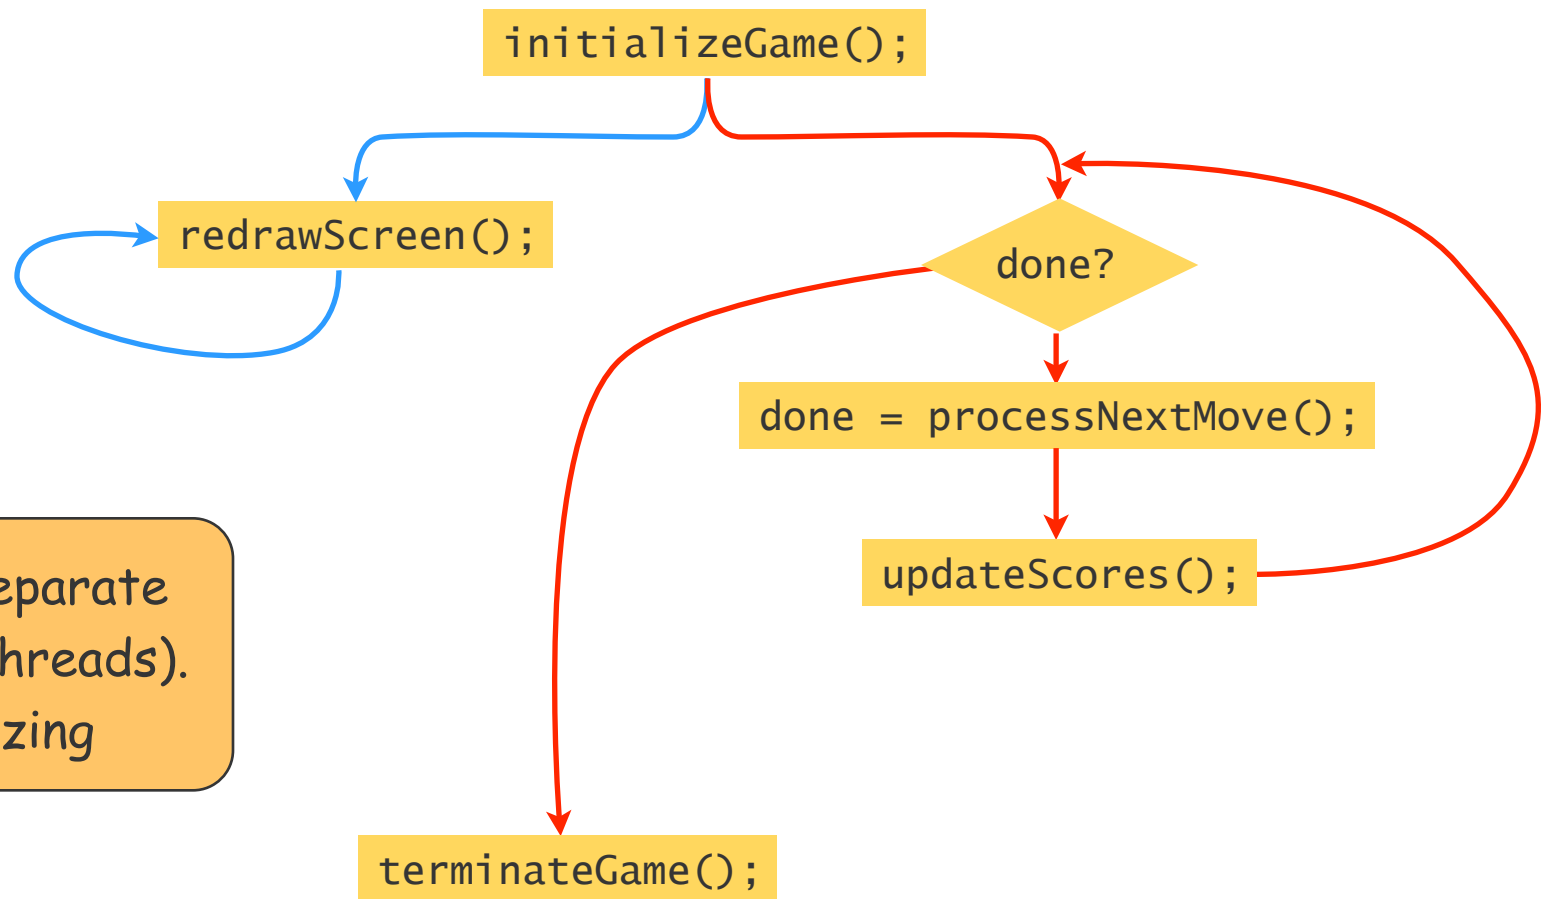
```
initializeGame();  
redrawScreen();  
boolean done=false;  
while(!done) {  
    done = processNextMove();  
    redrawScreen();  
    updateScores();  
}  
terminateGame();
```

Screen frozen  
while waiting  
for user input.





# [ Game: concurrent version ]




Note: separate  
tasks (threads).  
No freezing

# [ Array: sequential version ]

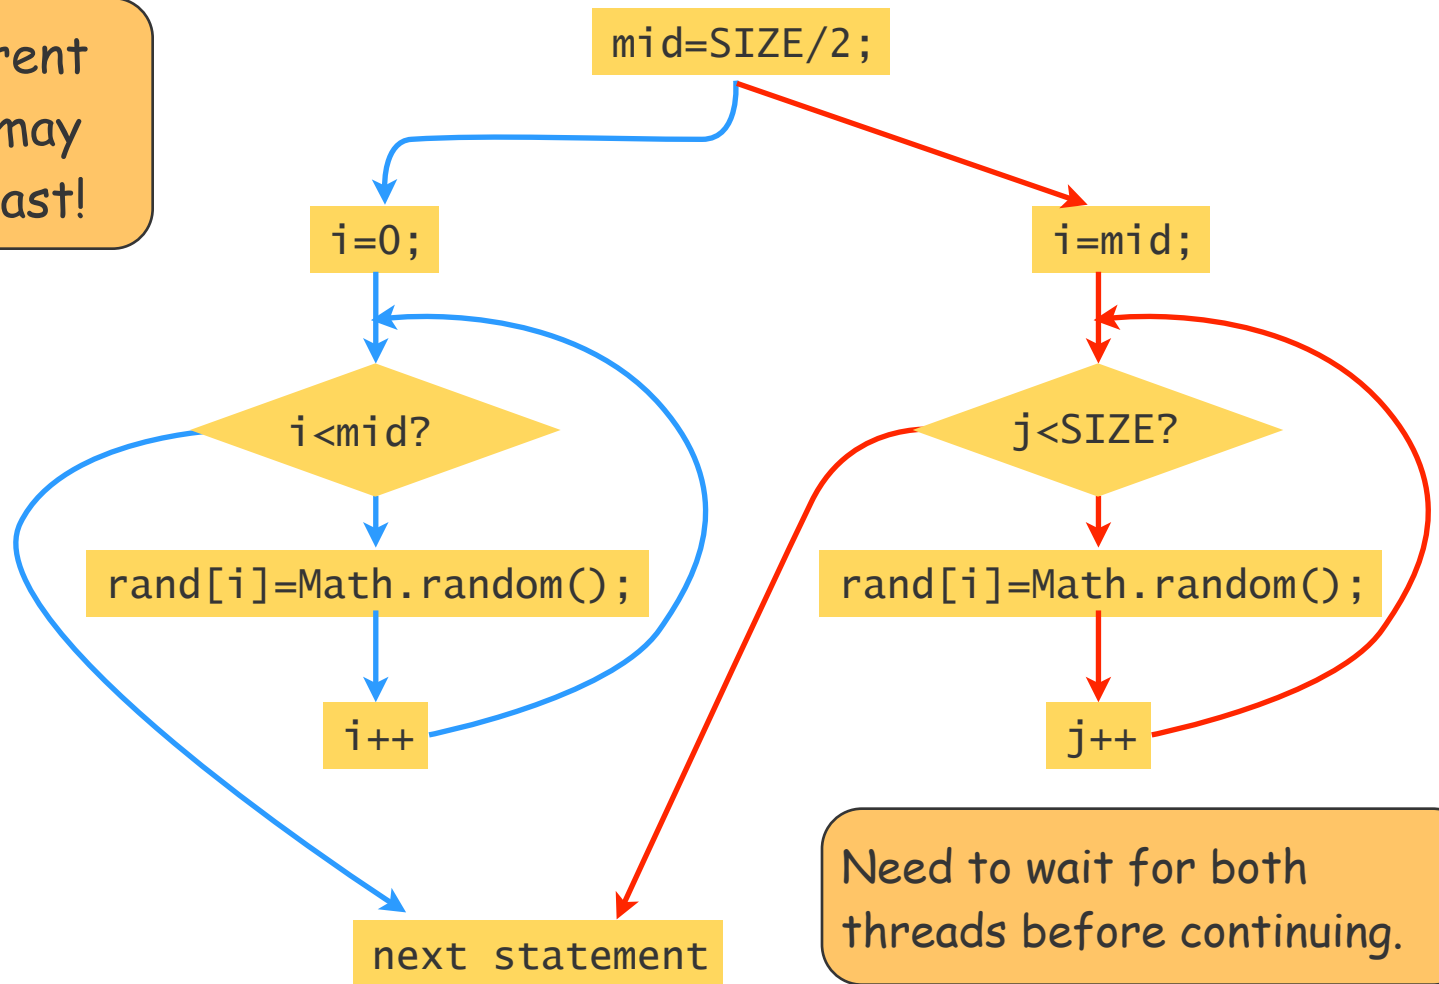
```
final int SIZE = 1000000;  
double[] rand = new double[SIZE];  
for(int i=0;i<SIZE;i++)  
    rand[i]= Math.random();
```

Only one thread --  
may take long time  
even though some  
CPUs are idle



# [ Array: concurrent version ]

With concurrent execution -- may be twice as fast!



# [ Motivation for concurrency ]

- Need for asynchrony
  - need to perform separate tasks
    - e.g., game, GUI examples
  - potential for increased speed with multiple CPUs/cores
    - e.g, matrix example
- Achieving these goals is not straightforward

# [ Sequential processing ]

- In a non-concurrent (sequential) program there is only one thread: the main thread.
- This thread executes the main method and then terminates.
  - the flow of control is determined by the main method
- With GUI elements,
  - a separate thread handle event: Event Dispatch Thread

# [ Concurrent processing ]

- How do we create a separate thread of execution?
- The **Thread** class provides a facility for creating separate threads.
  - Declare a class to be a descendant of Thread
  - Override the **run()** method to perform the necessary task(s) for the new thread
- When the **start()** method is called, the thread starts executing concurrently

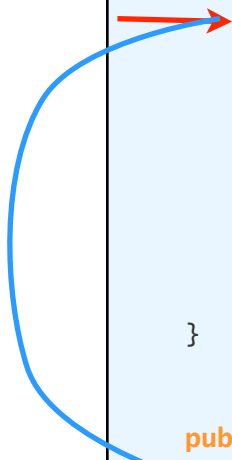
# Game: concurrent version

```
public class Game extends Thread{  
    public static void main(String[] args){  
        → Game game = new Game();  
        game.playGame();  
    }  
  
    public void playGame(){  
        boolean done=false;  
        initializeGame();  
        start(); ←  
        while(!done) {  
            done = processNextMove();  
            updateScores();  
        }  
        terminateGame();  
    }  
  
    public void run(){  
        while(true)  
            redrawScreen();  
    }  
}
```

NOTE!!

# Game: concurrent version

```
public class Game extends Thread{  
    public static void main(String[] args){  
        Game game = new Game();  
        game.playGame();  
    }  
  
    public void playGame(){  
        boolean done=false;  
        initializeGame();  
        start();  
        while(!done) {  
            done = processNextMove();  
            updateScores();  
        }  
        terminateGame();  
    }  
  
    public void run(){  
        while(true)  
            redrawScreen();  
    }  
}
```





# [ Array: concurrent version ]

```
public class ProcessArray {  
    static final int SIZE = 1000000;  
    static int[] data = new int[SIZE];  
  
    public static void main(String[] args){  
        → int mid = SIZE/2;  
        InitArray thread1 = new InitArray(0,mid, data);  
        InitArray thread2 = new InitArray(mid, SIZE, data);  
        thread1.start();  
        thread2.start();  
        ...  
    }  
}
```

```
public class InitArray extends Thread {  
    int start, end;  
    int array[];  
    public InitArray(int from, int to,  
        int[] array){  
        start = from;  
        end = to;  
        this.array = array;  
    }  
    public void run(){  
        → for(int i=start;i<end;i++)  
            array[i]= Math.random();  
    }  
}
```

NOTE!!

# [ Array: concurrent version

```
public class ProcessArray {  
    static final int SIZE = 1000000;  
    static int[] data = new int[SIZE];  
  
    public static void main(String[] args){  
        int mid = SIZE/2;  
        InitArray thread1 = new InitArray(0,mid, data);  
        InitArray thread2 = new InitArray(mid, SIZE, data);  
        thread1.start();  
        thread2.start();  
        ...  
    }  
}
```

```
public class InitArray extends Thread {  
    int start, end;  
    int array[];  
    public InitArray(int from, int to,  
        int[] array){  
        start = from;  
        end = to;  
        this.array = array;  
    }  
    public void run(){  
        for(int i=start;i<end;i++)  
            array[i]= Math.random();  
    }  
}
```

# [ Rejoining threads ]

- In the last example, it is necessary to wait for both threads to finish before moving on.
- This is achieved by calling the **join()** method
  - the thread that calls join is suspended until the thread on which it is called terminates.
  - this method can throw the (checked) **InterruptedException** so we should catch this exception

# Array: concurrent version 2

```
public class ProcessArray {  
    static final int SIZE = 1000000;  
    static int[] odd = new int[SIZE];  
  
    public static void main(String[] args){  
        int mid = SIZE/2;  
        InitArray thread1 = new InitArray(0,mid);  
        InitArray thread2 = new InitArray(mid, SIZE);  
        thread1.start();  
        thread2.start();  
  
        try{  
            thread1.join();  
            thread2.join();  
        } catch (InterruptedException e){  
            System.out.println("Error in thread");  
        }  
  
        . . .  
    }  
}
```

```
public class InitArray extends Thread {  
    int start, end;  
    int array[];  
    public InitArray(int from, int to,  
        int[] array){  
        start = from;  
        end = to;  
        this.array = array;  
    }  
    public void run(){  
        for(int i=start;i<end;i++){  
            array[i]= Math.random();  
        }  
    }  
}
```

# [The join() method]

- A call to the join method blocks (i.e., does not return) until the thread on which it is called terminates
  - returns from its run() method, or
  - propagates an exception from run()
- While being blocked, the calling thread may get interrupted which is why the join method throws the exception.
- Do not use the stop() method to stop a thread -- deprecated.

# [Speedup]

- Two key reasons for concurrency:
  - liveness (e.g., game keeps redrawing screen)
  - speedup (with more cores, programs run faster)
- Speedup can be measured using the System class methods:
  - public static long currentTimeMillis()
    - time elapsed since 1/1/1970 12:00am, in ms
  - public static long nanoTime()
    - current value of computer's timer in ns.

# [Game: concurrent version]

NOTE!!



```
public class Game extends Thread{  
    public static void main(String[]  
args){  
        Game game = new Game();  
        game.playGame();  
    }  
    public void playGame(){  
        boolean done=false;  
        initializeGame();  
        start();  
        while(!done) {  
            done = processNextMove();  
            updateScores();  
        }  
        terminateGame();  
    }  
    public void run(){  
        while(true)  
            redrawScreen();  
    }  
}
```

# [ Array: concurrent version ]

```
public class ProcessArray {  
    static final int SIZE = 1000000;  
    static int[] data = new int[SIZE];  
  
    public static void main(String[] args){  
        int mid = SIZE/2;  
        InitArray thread1 = new InitArray(0,mid, data);  
        InitArray thread2 = new InitArray(mid, SIZE, data);  
        thread1.start();  
        thread2.start();  
        ...  
    }  
}
```

```
public class InitArray extends Thread {  
    int start, end;  
    int array[];  
    public InitArray(int from, int to,  
        int[] array){  
        start = from;  
        end = to;  
        this.array = array;  
    }  
    public void run(){  
        for(int i=start;i<end;i++)  
            array[i]= Math.random();  
    }  
}
```

NOTE!!



# [ Creating Sub-Tasks ]

- To achieve concurrent processing, we need to divide a task into multiple pieces that can be assigned to concurrent threads.
- Two main approaches
  - Task decomposition
    - divide the type of work being performed
      - e.g., game example
  - Domain decomposition
    - divide the data on which the same task is performed
      - e.g. matrix initialization

# Array multiplication

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} = \begin{bmatrix} b_{11} & b_{12} & b_{13} & b_{14} \\ b_{21} & b_{22} & b_{23} & b_{24} \\ b_{31} & b_{32} & b_{33} & b_{34} \\ b_{41} & b_{42} & b_{43} & b_{44} \end{bmatrix} \times \begin{bmatrix} c_{11} & c_{12} & c_{13} & c_{14} \\ c_{21} & c_{22} & c_{23} & c_{24} \\ c_{31} & c_{32} & c_{33} & c_{34} \\ c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix}$$

$$a_{ij} = \sum_{k=1}^4 b_{ik} * c_{kj} = b_{i1} * c_{1j} + b_{i2} * c_{2j} + b_{i3} * c_{3j} + b_{i4} * c_{4j}$$

# [ Task sub-division ]

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} = \begin{bmatrix} b_{11} & b_{12} & b_{13} & b_{14} \\ b_{21} & b_{22} & b_{23} & b_{24} \\ b_{31} & b_{32} & b_{33} & b_{34} \\ b_{41} & b_{42} & b_{43} & b_{44} \end{bmatrix} \times \begin{bmatrix} c_{11} & c_{12} & c_{13} & c_{14} \\ c_{21} & c_{22} & c_{23} & c_{24} \\ c_{31} & c_{32} & c_{33} & c_{34} \\ c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix}$$

# Array Multiplication Example

```
public class ArrayMult {  
    public static void main(String[] args){  
        final int M=4, N=4, K=4;  
        int[][] a = new int[M][N];  
        int[][] b = new int[M][K];  
        int[][] c = new int[K][N];  
        . . . //initialize values;  
        ArrayMult mult1 = new ArrayMult(0,M/2,a,b,c);  
        ArrayMult mult2 = new ArrayMult(M/2,M,a,b,c);  
        mult1.start();  
        mult2.start();  
        try{  
            mult1.join();  
            mult2.join();  
        } catch (InterruptedException e) {  
            System.out.println("Unexpected  
Interrupt");  
        }  
    }  
}
```

```
public class ArrayMult extends Thread {  
    int start, end;  
    int[][] a,b,c;  
    public ArrayMult(int from, int to, int  
        [][] a, int[][] b, int[][] c){  
        start = from;  
        end = to;  
        this.a = a;  
        this.b = b;  
        this.c = c;  
    }  
    public void run(){  
        for(int i=start;i<end;i++)  
            for(int j=0;j<a[0].length;j++){  
                a[i][j] = 0;  
                for(int k=0;k<b.length;k++){  
                    a[i][j]+=b[i][k]*c[k][j];  
                }  
            }  
    }  
}
```

# Task sub-division

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} = \begin{bmatrix} b_{11} & b_{12} & b_{13} & b_{14} \\ b_{21} & b_{22} & b_{23} & b_{24} \\ b_{31} & b_{32} & b_{33} & b_{34} \\ b_{41} & b_{42} & b_{43} & b_{44} \end{bmatrix} \times \begin{bmatrix} c_{11} & c_{12} & c_{13} & c_{14} \\ c_{21} & c_{22} & c_{23} & c_{24} \\ c_{31} & c_{32} & c_{33} & c_{34} \\ c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix}$$

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} = \begin{bmatrix} b_{11} & b_{12} & b_{13} & b_{14} \\ b_{21} & b_{22} & b_{23} & b_{24} \\ b_{31} & b_{32} & b_{33} & b_{34} \\ b_{41} & b_{42} & b_{43} & b_{44} \end{bmatrix} \times \begin{bmatrix} c_{11} & c_{12} & c_{13} & c_{14} \\ c_{21} & c_{22} & c_{23} & c_{24} \\ c_{31} & c_{32} & c_{33} & c_{34} \\ c_{41} & c_{42} & c_{43} & c_{44} \end{bmatrix}$$

```

public class ArrayMult {
    public static void main(String[] args){
        final int NUM_THREADS=5;
        ArrayMult[] threads = new ArrayMult[NUM_THREADS];
        . . .
        int subsetSize=(int) Math.ceil(a.length /(float)NUM_THREADS);
        int startRow = 0;
        for(int i=0;i<NUM_THREADS;i++){
            threads[i]=new ArrayMult(startRow, Math.min(startRow+subsetSize,a.length), a,
            b, c);
            threads[i].start();
            start+=subsetSize;
        }
        try{
            for(int i=0;i<NUM_THREADS;i++)
                threads[i].join();
        } catch (InterruptedException e) {
            System.out.println("Unexpected Interrupt");
        }
    }
}

```

# [Processes]

- Modern operating systems support multi-tasking
  - painting the screen, listening to the keyboard, printing, running several programs, ...
- Even with a single core multiple tasks are concurrently running
- Achieved by sharing the processor among multiple processes
  - the CPU runs a little of each process in turn
  - this is called *process scheduling*

# [Threads]

- A process often corresponds to a program
  - Browser, editor, ...
- Modern processes often have multiple threads of execution.
- Roughly,
  - different processes are largely independent of each other;
  - different threads of the same process often share the same memory space.



# [ Thread Scheduling ]

- Within a single thread, instructions are processed one at a time.
- However, different threads can run at different times/rates.
- When a thread runs is determined by many factors:
  - Java implementation;
  - Operating system
  - Instructions being executed
  - ....

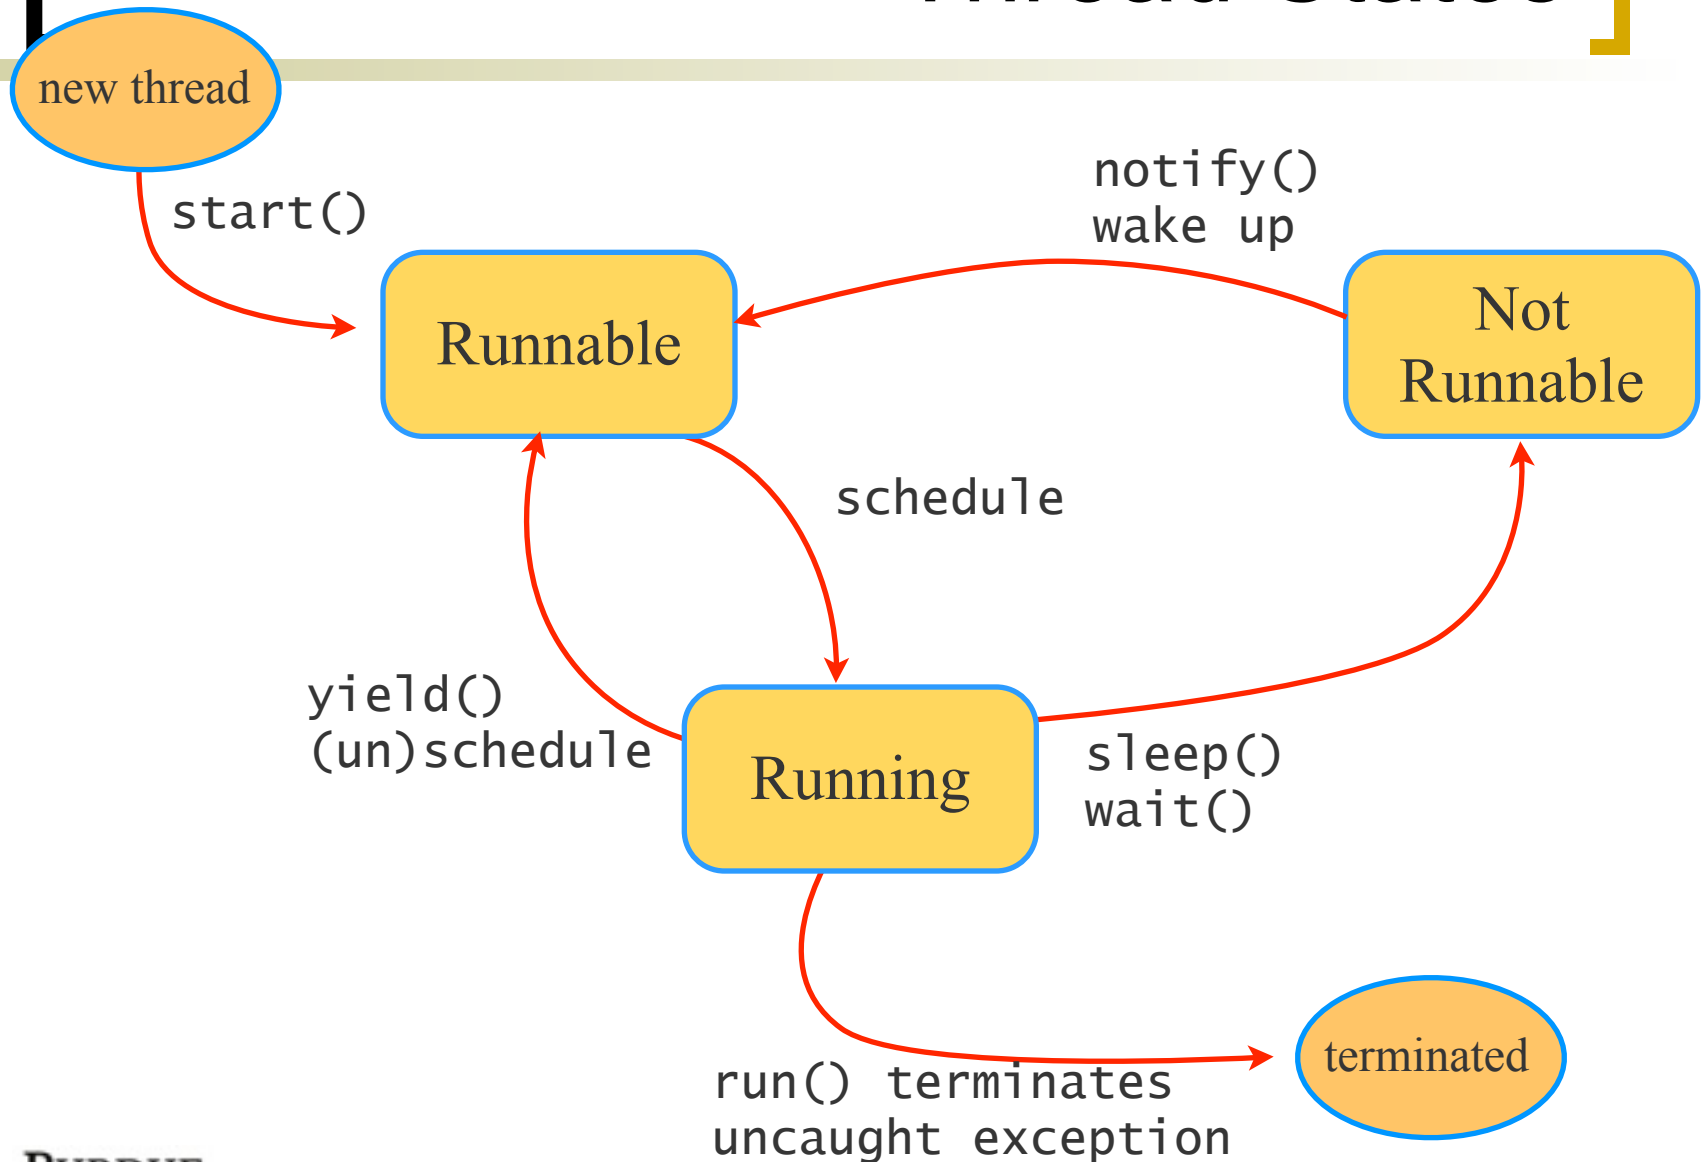
# [ Non-determinism ]

- An important property of threads is that it is not possible to know exactly when a given thread will be scheduled
  - cannot assume anything about relative ordering between threads (more later)
- Order of concurrent threads (and consequently the result of the output) may change from run to run!
- Programmer must anticipate all possible orderings and protect against possible errors.

# [Controlling thread scheduling]

- As a programmer we have several mechanism available:
  - sleep()
    - thread cannot be scheduled for some time
  - yield()
    - voluntarily give up your turn for the CPU
  - wait()
    - wait for some condition to be true
  - Priority
    - Each thread has a priority. Can set priorities for threads we create (with some limitations).

# Thread States



# [ Thread scheduling ]

- At any given time there may be a number of threads that are runnable
  - each has a priority
  - usually the same as the creating thread's priority
- Periodically, the OS schedules one of the threads with the highest priority for some time.

# [ Synchronization example ]

- Say we want to try to control the relative ordering of two threads:
  - thread1 prints: “Left, Left, Left” then “Left”
  - thread2 prints: “Right”
- Suppose we want to ensure the following output:
  - “Left, Left, Left, Right, Left” multiple times.
  - How can we ensure that the timing of the threads ensures this output?
  - I.e., how to avoid non-determinism?

# [ Attempt 1: using sleep() ]

```
public class LeftThread extends Thread {  
    public void run(){  
        for(int i=0;i<10;i++){  
            System.out.print("Left ");  
            System.out.print("Left ");  
            System.out.print("Left ");  
            try {  
                Thread.sleep(50);  
            } catch (InterruptedException e){  
                e.printStackTrace();  
            }  
            System.out.println("Left ");  
        }  
    }  
}
```

```
public class RightThread extends Thread {  
    public void run(){  
        try {  
            Thread.sleep(10);  
        } catch (InterruptedException e) {  
            e.printStackTrace();  
        }  
        for(int i=0;i<10;i++){  
            System.out.print("Right ");  
            try {  
                Thread.sleep(50);  
            } catch (InterruptedException e) {  
                e.printStackTrace();  
            }  
        }  
    }  
}
```

# [Problems with sleep()]

- Doesn't work
  - There is no guarantee that with the sleeping we will get synchronized each time
  - With enough chances, will get out of sync
- There may be unnecessary waiting
- Hard to tune the sleep times



# [ Attempt 2: using yield() ]

```
public class LeftThread extends Thread {  
    public void run(){  
        for(int i=0;i<10;i++){  
            System.out.print("Left ");  
            System.out.print("Left ");  
            System.out.print("Left ");  
            Thread.yield();  
            System.out.println("Left ");  
        }  
    }  
}
```

```
public class RightThread extends Thread {  
    public void run(){  
        Thread.yield();  
        for(int i=0;i<10;i++){  
            System.out.print("Right ");  
            Thread.yield();  
        }  
    }  
}
```

# [Problems with yield()]

- Assumes that the yield() calls will give control to the other thread
  - many threads may be running on the machine
  - can cause unexpected switches between our 2 threads
- With multiple cores, each will be running on a separate core -- yielding does not provide anything!
- Also doesn't work

# Polling

```
public class LeftThread extends Thread {
    private RightThread right;
    private boolean done = false;
    public void setRight(RightThread right){
        this.right=right;
    }
    public void run(){
        for(int i=0;i<reps;i++){
            System.out.print("Left ");
            System.out.print("Left ");
            System.out.print("Left ");
            done = true;
            while(!right.isDone());
            right.setDone(false);
            Thread.yield();
            System.out.println("Left ");
        }
    }
    public boolean isDone() {return done;}
    public void setDone(boolean value) {done = value;}
}
```

```
public class RightThread extends Thread {
    private LeftThread left;
    private boolean done = false;
    public void setLeft(LeftThread left){
        this.left = left;
    }
    public void run(){
        for(int i=0;i<reps;i++){
            while(!left.isDone()) ;
            left.setDone(false);
            System.out.print("Right ");
            done = true;
        }
    }
    public boolean isDone() {return done;}
    public void setDone(boolean value) {
        done = value;
    }
}
```

# [ Polling solution? ]

- This works
  - always produces correct output.
- However,
  - No real concurrency!
  - Only one thread running at a time.
  - Busy waiting (wastes resources)
- Technicality:
  - should ensure that *done* variables are visible to the other thread immediately: use the **volatile** modifier.

# [ Shared Memory Architecture ]

- Two common approaches to concurrent programs:
  - message passing
  - shared memory
- Java uses shared memory
  - multiple threads of the same application (program) essentially have access to the same memory space (i.e., variables)
  - memory on each core/CPU is not shared
  - can lead to delays in visibility of modifications (use volatile to avoid these if multiple threads will modify the same variable).

# [Concurrency is tricky]

- Writing concurrent programs that work as expected can be tricky
- Need to deal with
  - non-determinism of scheduling
  - ensuring access to shared data is correct (see slides on Synchronization)
- Achieving speed up is not always easy

# [ Examples ]

- Factorization of a large integer
  - need to find the two prime factors of a large integer value
  - divide the task by domain decomposition
- Array summation
  - compute the sum of the sine of all values of a large array
  - divide by domain decomposition
  - need to synchronize after sub-tasks are done

# Factorization

```
public class FactorThread extends Thread {
    private long lower;
    private long upper;
    public static final int THREADS = 4;
    public static final long NUMBER = 59984005171248659L;
    public FactorThread(long lower, long upper){
        this.lower = lower;
        this.upper = upper;
    }
    public void run(){
        if(lower%2==0)
            lower++;
        while(lower<upper) {
            if(NUMBER%lower == 0) {
                System.out.println("Security Code: " + (lower + NUMBER/lower));
                return;
            }
            lower += 2;
        }
    }
    public static void main (String[] args ) {...}
}
```



# Factorization (main)

```
public static void main(String[] args){
    FactorThread[] threads = new FactorThread[THREADS];
    long root = (long)Math.sqrt(NUMBER);
    long start = 3;
    long numbers = (long)Math.ceil((root-2)/(float)THREADS);

    for(int i=0;i<THREADS;i++){
        threads[i] = new FactorThread(start, Math.min(start
+numbers, root+1));
        threads[i].start();
        start+=numbers;
    }
    try{
        for(int i=0;i<THREADS;i++)
            threads[i].join();
    } catch (InterruptedException e){
        e.printStackTrace();
    }
}
```

# Matrix Sum

```
import java.util.Random ;
public class SumThread extends Thread {
    private static double [] data ;
    private static SumThread [] threads ;
    private double sum = 0;
    private int lower, upper, index ;
    public static final int SIZE = 1000000;
    public static final int THREADS = 8;
    public SumThread (int lower, int upper, int index) {
        this.lower = lower;
        this.upper = upper;
        this.index = index;
    }
    public double getSum () { return sum ; }
    public void run () { //next slide }
    public static void main ( String [] args ) { // next slide }
}
```

# Matrix Sum (contd.)

```
public void run () {  
    for ( int i = lower ; i < upper ; i++ )  
        sum += Math.sin( data [i]);  
    int power = 2;  
    int neighbor;  
    while (index % power == 0 && power < THREADS) {  
        neighbor = index + power / 2;  
        try { threads [ neighbor ].join (); }  
        catch ( InterruptedException e ) {  
            e.printStackTrace ();  
        }  
        sum += threads [ neighbor ].getSum ();  
        power *= 2;  
    }  
}
```

# Matrix (contd.)

```
public static void main ( String [] args ) {  
    data = new double [ SIZE ];  
    Random random = new Random ();  
    int start = 0;  
    for ( int i = 0; i < SIZE ; i++ )  
        data [i] = random . nextDouble ();  
    threads = new SumThread [ THREADS ];  
    int range = ( int) Math . ceil ( data . length / ( float ) THREADS );  
    for ( int i = 0; i < THREADS ; i++ ) {  
        threads [i] = new SumThread ( start, Math . min( start + range , SIZE ), i );  
        threads [i]. start ();  
        start += range ;  
    }  
    try { threads [0]. join (); }  
    catch ( InterruptedException e ) {  
        e. printStackTrace ();  
    }  
    System .out. println ( "Sum: " + threads [0]. getSum ());  
}
```